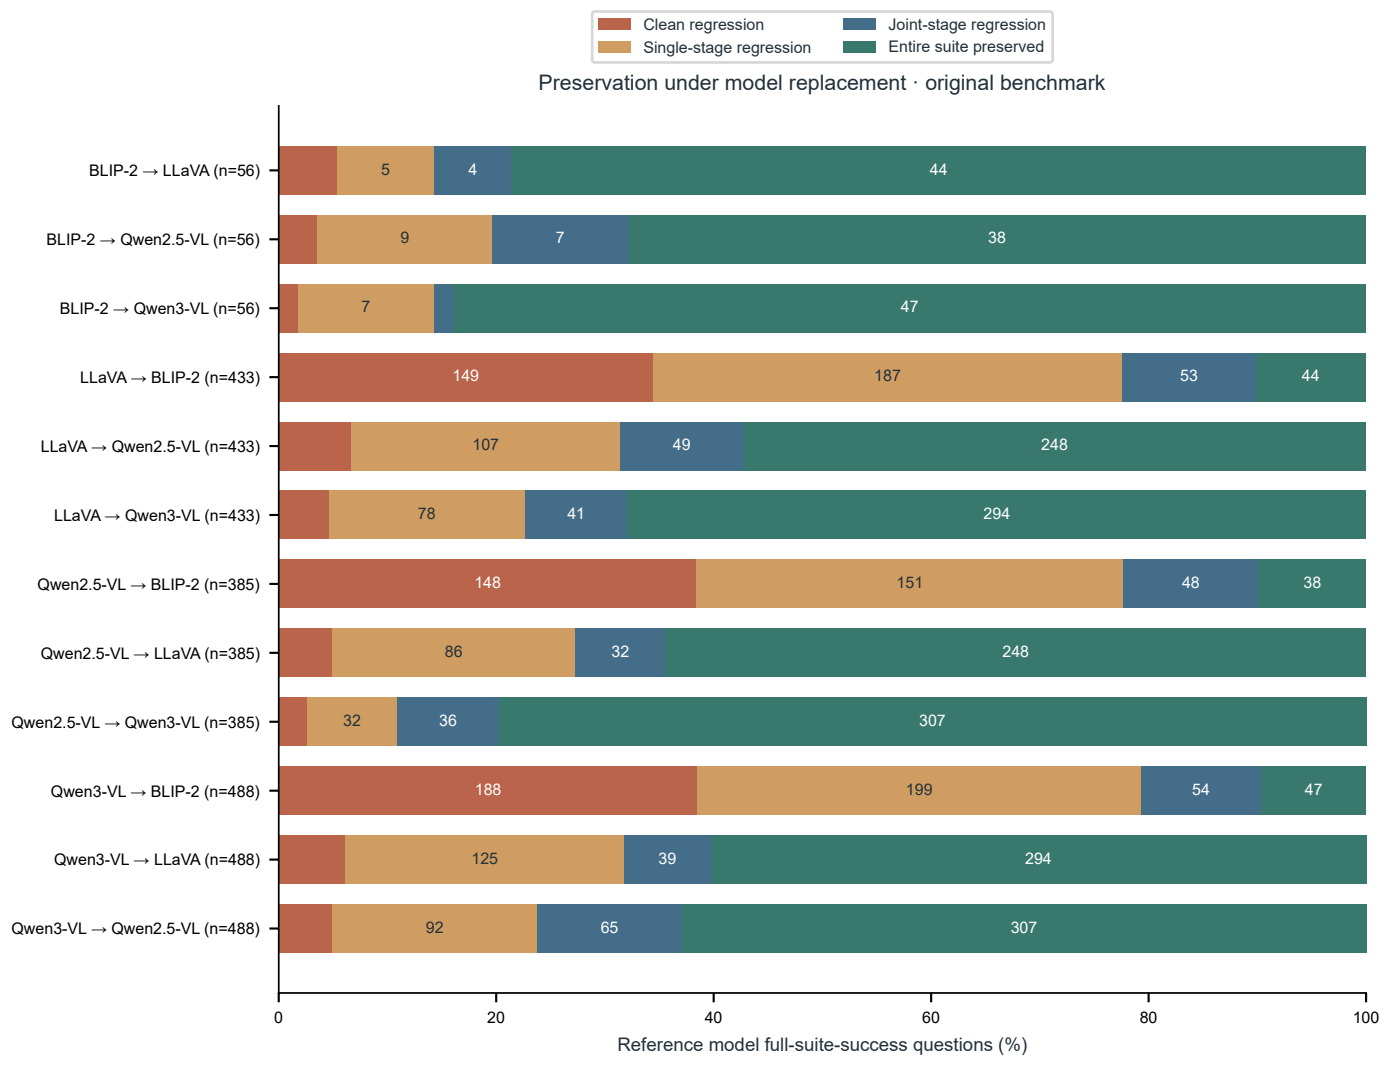
Clean regression
Single-stage regression
Joint-stage regression
Entire suite preserved
Preservation under model replacement · original benchmark
BLIP-2 → LLaVA (n=56)
5
4
44
BLIP-2 → Qwen2.5-VL (n=56)
9
7
38
BLIP-2 → Qwen3-VL (n=56)
7
47
LLaVA → BLIP-2 (n=433)
149
187
53
44
LLaVA → Qwen2.5-VL (n=433)
107
49
248
LLaVA → Qwen3-VL (n=433)
78
41
294
Qwen2.5-VL → BLIP-2 (n=385)
148
151
48
38
Qwen2.5-VL → LLaVA (n=385)
86
32
248
Qwen2.5-VL → Qwen3-VL (n=385)
32
36
307
Qwen3-VL → BLIP-2 (n=488)
188
199
54
47
Qwen3-VL → LLaVA (n=488)
125
39
294
Qwen3-VL → Qwen2.5-VL (n=488)
92
65
307
0
20
40
60
80
100
Reference model full-suite-success questions (%)